

VEHICLE SELECTION USING RIDGE AND LASSO

In [5]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [6]:

```
df=pd.read_csv(r"C:\Users\RAMADEVI SURIPAKA\Downloads\fiat500_VehicleSelection_Dataset (
df
```

Out[6]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat
0	1	lounge	51	882	25000	1	44.907242
1	2	pop	51	1186	32500	1	45.666359
2	3	sport	74	4658	142228	1	45.503300
3	4	lounge	51	2739	160000	1	40.633171
4	5	pop	73	3074	106880	1	41.903221
...
1533	1534	sport	51	3712	115280	1	45.069679
1534	1535	lounge	74	3835	112000	1	45.845692
1535	1536	pop	51	2223	60457	1	45.481541
1536	1537	lounge	51	2557	80750	1	45.000702
1537	1538	pop	51	1766	54276	1	40.323410

1538 rows × 9 columns



In [7]:

```
df.head()
```

Out[7]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	p
0	1	lounge	51	882	25000	1	44.907242	8.611560	1
1	2	pop	51	1186	32500	1	45.666359	12.241890	1
2	3	sport	74	4658	142228	1	45.503300	11.417840	1
3	4	lounge	51	2739	160000	1	40.633171	17.634609	1
4	5	pop	73	3074	106880	1	41.903221	12.495650	1

In [8]:

```
df.tail()
```

Out[8]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	p
1533	1534	sport	51	3712	115280	1	45.069679	7.7041	1
1534	1535	lounge	74	3835	112000	1	45.845692	8.6661	1
1535	1536	pop	51	2223	60457	1	45.481541	9.4131	1
1536	1537	lounge	51	2557	80750	1	45.000702	7.6821	1
1537	1538	pop	51	1766	54276	1	40.323410	17.5681	1

In [9]:

```
df.shape
```

Out[9]:

(1538, 9)

In [10]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    1538 non-null   int64
1   model                 1538 non-null   object
2   engine_power          1538 non-null   int64
3   age_in_days           1538 non-null   int64
4   km                    1538 non-null   int64
5   previous_owners       1538 non-null   int64
6   lat                   1538 non-null   float64
7   lon                   1538 non-null   float64
8   price                 1538 non-null   int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```

In [11]:

```
df.isna().any()
```

Out[11]:

```
ID                False
model              False
engine_power       False
age_in_days        False
km                 False
previous_owners    False
lat                False
lon                False
price              False
dtype: bool
```

In [12]:


```
df.describe()
```

Out[12]:

	ID	engine_power	age_in_days	km	previous_owners	li
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.54136
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.13351
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.85583
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.80295
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.39405
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.46796
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.79561

In [13]:

```
sns.lmplot(x='km',y='price',data=df,order=2,ci=None)  
plt.show()
```



In [14]:

```
x=np.array(df['km']).reshape(-1,1)  
y=np.array(df['price']).reshape(-1,1)
```

In [15]:

```
df.dropna(inplace=True)
```


In [16]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)  
#splitting data into train and test  
regr=LinearRegression()  
regr.fit(x_train,y_train)  
print(regr.score(x_test,y_test))
```

0.7671660418686099


In [17]:

```
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



In [18]:

```
df500=df[:][:500]
sns.lmplot(x="km",y="price",data=df500,order=1,ci=None)
plt.show()
```



In [19]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

In [20]:

```
#train model
model=LinearRegression()
model.fit(x_train,y_train)
#Evaluation the model on the test set
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 score:",r2)
```

R2 score: 0.7671660418686099

RIDGE AND LASSO

In [21]:

```
from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

In [22]:

```
features=df.columns[0:1]
target=df.columns[-1]
```

In [27]:

```
converter={"model":{"sport":1,"lounge":2,"pop":3}}
df=df.replace(converter)
df
```

Out[27]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	
0	1	2	51	882	25000	1	44.907242	8.6111
1	2	3	51	1186	32500	1	45.666359	12.2411
2	3	1	74	4658	142228	1	45.503300	11.4171
3	4	2	51	2739	160000	1	40.633171	17.6341
4	5	3	73	3074	106880	1	41.903221	12.4951
...
1533	1534	1	51	3712	115280	1	45.069679	7.7041
1534	1535	2	74	3835	112000	1	45.845692	8.6661
1535	1536	3	51	2223	60457	1	45.481541	9.4131
1536	1537	2	51	2557	80750	1	45.000702	7.6821
1537	1538	3	51	1766	54276	1	40.323410	17.5681

1538 rows × 9 columns




In [28]:

```
plt.figure(figsize = (10, 10))
sns.heatmap(df.corr(), annot = True)
```

Out[28]:

<Axes: >



In [24]:

```
features=df.columns[0:1]
target=df.columns[-1]
```


In [44]:

```
#X and y values
X = df[features].values
y = df[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=1)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of X_train is (1153, 1)
The dimension of X_test is (385, 1)

In [45]:

```
lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.00310286926477088
The test score for lr model is -0.008405634316406507

In [46]:


```
#Ridge Regression Model
ridgeReg = Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test scorefor ridge regression
train_score_ridge = ridgeReg.score(X_train, y_train)
test_score_ridge = ridgeReg.score(X_test, y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.0031026398591535997
The test score for ridge model is -0.008307809466001403

In [47]:

```
plt.figure(figsize=(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color='green')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker="o",markersize=7,color='green')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```



In [48]:

```
#Lasso regression model
print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.003075838461310987


The test score for ls model is -0.007367578602064606

In [49]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[49]:

<Axes: >



In [50]:


```
#Using the linear CV model
from sklearn.linear_model import LassoCV
#Lasso Cross validation
lasso_cv = LassoCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
#score
print(lasso_cv.score(X_train, y_train))
print(lasso_cv.score(X_test, y_test))
```

0.0031025989567363688

-0.008299466692577973

In [54]:

```
#plot size
plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red')
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',)
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green')
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```



In [52]:

```
#Using the linear CV model
from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)))
print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))
```

The train score for ridge model is 0.0031026398591535997

The train score for ridge model is -0.008307809466002958

Type *Markdown* and LaTeX: α^2

In []:

In []:

In []: